

# AI ASSISTED CODING

Name: N.Prudhvi

Htno: 2403A510G7

Batch: 06

## Lab-Assignment: 17.1

### Task-1:

#### Prompt:

Write Python code to clean social media posts by removing stopwords, punctuation, and special symbols from text, handling missing values in likes and shares using the median, converting timestamps to datetime, extracting hour and weekday features, and removing duplicate or spam posts.

#### Code:

```
1 import pandas as pd
2 import re
3 from typing import List
4
5 # A basic list of English stopwords. For more comprehensive lists, consider u
6 STOPWORDS = set(
7     [
8         "i",
9         "me",
10        "my",
11        "myself",
12        "we",
13        "our",
14        "ours",
15        "ourselves",
16        "you",
17        "your",
18        "yours",
19        "he",
20        "him",
20        "him",
21        "his",
22        "she",
23        "her",
24        "it",
25        "its",
26        "they",
27        "them",
28        "their",
29        "what",
30        "which",
31        "who",
32        "whom",
33        "this",
34        "that",
35        "these",
36        "those",
37        "am",
38        "is",
39        "are",
```

```

130
131 v def clean_text(text: str) -> str:
132     """Removes stopwords, punctuation, and special symbols from a text string.
133 v     if not isinstance(text, str):
134         return ""
135     # Convert to lowercase
136     text = text.lower()
137     # Remove URLs, mentions, and hashtags
138     text = re.sub(r"http\S+|www\S+|@\w+|#\w+", "", text)
139     # Remove punctuation and special symbols
140     text = re.sub(r"[^a-z\s]", "", text)
141     # Remove stopwords
142     words = [word for word in text.split() if word not in STOPWORDS]
143     return " ".join(words)
144
145
146 v def clean_social_media_data(df: pd.DataFrame) -> pd.DataFrame:
147 v     """
148     Cleans a raw social media dataset by handling missing values, cleaning text

```

```

148     Cleans a raw social media dataset by handling missing values, cleaning text
149     processing timestamps, and removing duplicates/spam.
150     """
151     # Make a copy to avoid modifying the original DataFrame
152     cleaned_df = df.copy()
153
154     # 1. Handle missing values in likes and shares by filling with the median
155     for col in ["likes", "shares"]:
156         if col in cleaned_df.columns:
157             # Coerce to numeric, turning non-numeric into NaN
158             cleaned_df[col] = pd.to_numeric(cleaned_df[col], errors="coerce")
159             median_val = cleaned_df[col].median()
160             cleaned_df[col].fillna(median_val, inplace=True)
161             cleaned_df[col] = cleaned_df[col].astype(int) # Convert to integer
162
163     # 2. Remove stopwords, punctuation, and special symbols
164     cleaned_df["cleaned_text"] = cleaned_df["post_text"].apply(clean_text)
165

```

```

146 def clean_social_media_data(df: pd.DataFrame) -> pd.DataFrame:
170
171     # 4. Detect and remove spam/duplicate posts
172     # Simple spam detection based on keywords
173     spam_keywords = ["buy now", "free money", "click here"]
174     spam_pattern = "|".join(spam_keywords)
175     cleaned_df["is_spam"] = cleaned_df["post_text"].str.contains(
176         spam_pattern, case=False, na=False
177     )
178
179     # Remove duplicates based on the cleaned text content
180     cleaned_df.drop_duplicates(subset=["cleaned_text"], keep="first", inplace=
181
182     # Filter out posts flagged as spam
183     cleaned_df = cleaned_df[~cleaned_df["is_spam"]]
184
185     # Final cleanup: select and reorder columns
186     final_cols = [
187         "post_id",
188         "timestamp",
189         "cleaned_text",
190         "likes",
191         "shares",
192         "hour_of_day",
193         "day_of_week",
194     ]
195     # Ensure all expected columns (variable) final_cols: list[str]
196     final_cols = [col for col in final_cols if col in cleaned_df.columns]
197     cleaned_df = cleaned_df[final_cols]
198
199     return cleaned_df.reset_index(drop=True)
200
201
202 # --- Example Usage ---
203 if __name__ == "__main__":
204     # Sample raw social media data
205     raw_data = {
206         "post_id": [1, 2, 3, 4, 5, 6, 7, 8],
207         "timestamp": [

```

```

208         "2023-10-26 08:30:00",
209         "2023-10-26 09:15:00",
210         "2023-10-26 10:00:00",
211         "2023-10-26 11:00:00",
212         "2023-10-26 12:45:00",
213         "2023-10-26 14:20:00",
214         "2023-10-27 15:00:00",
215         "2023-10-27 16:00:00",
216     ],
217     "post_text": [
218         "Just had an amazing breakfast! * #foodie",
219         "This is a great article on AI: http://example.com/ai",
220         "Feeling tired today... need coffee ☕",
221         "!!! BUY NOW, limited offer !!!", # Spam post
222         "Just had an amazing breakfast! * #foodie", # Duplicate post
223         "What a game last night! Simply incredible.",
224         "Working on a new project. It is very exciting.",
225         "Another spam post with free money", # Spam post
226     ],
227     "likes": [150, 200, 75, 10, 120, 300, None, 5], # Includes a missing
    ],
    "likes": [150, 200, 75, 10, 120, 300, None, 5], # Includes a missing
    "shares": [20, 45, None, 1, 15, 80, 25, 0], # Includes a missing value
}
raw_df = pd.DataFrame(raw_data)

print("--- Original Raw Dataset ---")
print(raw_df)
print("\n" + "=" * 50 + "\n")

# Clean the dataset
cleaned_df = clean_social_media_data(raw_df)

print("--- Cleaned Dataset for Analysis ---")
print(cleaned_df)

```

**Output:**

```
cleaned_df[col].fillna(median_val, inplace=True)
--- Cleaned Dataset for Analysis ---
  post_id      timestamp  ... hour_of_day  day_of_week
0      1  2023-10-26 08:30:00  ...         8    Thursday
1      2  2023-10-26 09:15:00  ...         9    Thursday
2      3  2023-10-26 10:00:00  ...        10    Thursday
3      6  2023-10-26 14:20:00  ...        14    Thursday
4      7  2023-10-27 15:00:00  ...        15     Friday

[5 rows x 7 columns]
```

```
C:\Users\venub\OneDrive\Desktop\AIAC_Lab\Lab-17>python 17_1.py
--- Original Raw Dataset ---
  post_id      timestamp  ...  likes  shares
0      1  2023-10-26 08:30:00  ...  150.0   20.0
1      2  2023-10-26 09:15:00  ...  200.0   45.0
2      3  2023-10-26 10:00:00  ...   75.0    NaN
3      4  2023-10-26 11:00:00  ...   10.0    1.0
4      5  2023-10-26 12:45:00  ...  120.0   15.0
5      6  2023-10-26 14:20:00  ...  300.0   80.0
6      7  2023-10-27 15:00:00  ...    NaN   25.0
7      8  2023-10-27 16:00:00  ...    5.0    0.0

[8 rows x 5 columns]
```

[Live Share](#)

## Observation:

This task focuses on cleaning and preprocessing social media data to make it suitable for analysis. It involves removing noise from the text, such as stopwords, punctuation, URLs, mentions, and hashtags.

Additionally, it handles missing values in numeric columns like likes and shares by replacing them with median values, and ensures these columns are of integer type. Time-related features, such as the hour of day and day of the week, are extracted from timestamps to allow for temporal analysis. The task also identifies and removes duplicate posts as well as spam posts containing keywords like “buy now” or “free money,” resulting in a cleaner and more reliable dataset for downstream tasks.

## Task-2:

### Prompt:

Write Python code to preprocess stock market data by handling missing values in closing\_price and volume, creating 1-day and 7-day lag return features, applying log-scaling to volume, and detecting outliers in closing\_price using the IQR method.

### Code:

```
1  import pandas as pd
2  import numpy as np
3
4
5  def preprocess_stock_data(df: pd.DataFrame) -> pd.DataFrame:
6      """
7      Preprocess a stock market dataset by handling missing values,
8      creating lag features, normalizing volume, and detecting outliers.
9      """
10     df = df.copy()
11
12     # 1. Handle missing values in 'closing_price' and 'volume'
13     for col in ["closing_price", "volume"]:
14         if col in df.columns:
15             df[col] = pd.to_numeric(df[col], errors="coerce")
16             df[col].fillna(method="ffill", inplace=True) # Forward fill
17             df[col].fillna(
18                 method="bfill", inplace=True
19             ) # Backward fill if first row is NaN
20
21     # 2. Create lag features (1-day and 7-day returns)
22     df["return_1d"] = df["closing_price"].pct_change(1)
23     df["return_7d"] = df["closing_price"].pct_change(7)
24
25     # 3. Normalize volume using log scaling
26     df["volume_log"] = df["volume"].apply(
27         lambda x: np.log1p(x)
28     ) # Log(1 + x) to avoid Log(0)
29
30     # 4. Detect outliers in 'closing_price' using IQR method
```

```

31     Q1 = df["closing_price"].quantile(0.25)
32     Q3 = df["closing_price"].quantile(0.75)
33     IQR = Q3 - Q1
34     lower_bound = Q1 - 1.5 * IQR
35     upper_bound = Q3 + 1.5 * IQR
36     df["is_outlier"] = (df["closing_price"] < lower_bound) | (
37         df["closing_price"] > upper_bound
38     )
39
40     # Reset index
41     return df.reset_index(drop=True)
42
43
44 # --- Example Usage ---
45 if __name__ == "__main__":
46     # Sample stock dataset
47     data = {
48         "date": pd.date_range(start="2023-10-20", periods=10, freq="D"),
49         "closing_price": [100, 102, 101, None, 105, 107, 106, 108, 110, 109],
50         "volume": [5000, 5200, None, 5400, 5600, None, 5800, 6000, 6200, 6400]
51     }
52
53     df = pd.DataFrame(data)
54     print("--- Original Stock Data ---")
55     print(df)
56
57     processed_df = preprocess_stock_data(df)
58     print("\n--- Preprocessed Stock Data ---")
59     print(processed_df)

```

## Output:

```

C:\Users\venub\OneDrive\Desktop\AIAC_Lab\Lab-17>python 17_1.py
--- Original Stock Data ---
   date   closing_price  volume
0 2023-10-20         100.0   5000.0
1 2023-10-21         102.0   5200.0
2 2023-10-22         101.0     NaN
3 2023-10-23          NaN   5400.0
4 2023-10-24         105.0   5600.0
5 2023-10-25         107.0     NaN
6 2023-10-26         106.0   5800.0
7 2023-10-27         108.0   6000.0
8 2023-10-28         110.0   6200.0
9 2023-10-29         109.0   6400.0

```



--- Preprocessed Stock Data ---							
	date	closing_price	volume	return_1d	return_7d	volume_log	is_outlier
0	2023-10-20	100.0	5000.0	NaN	NaN	8.517393	False
1	2023-10-21	102.0	5200.0	0.020000	NaN	8.556606	False
2	2023-10-22	101.0	5200.0	-0.009804	NaN	8.556606	False
3	2023-10-23	101.0	5400.0	0.000000	NaN	8.594339	False
4	2023-10-24	105.0	5600.0	0.039604	NaN	8.630700	False
0	2023-10-20	100.0	5000.0	NaN	NaN	8.517393	False
1	2023-10-21	102.0	5200.0	0.020000	NaN	8.556606	False
2	2023-10-22	101.0	5200.0	-0.009804	NaN	8.556606	False
3	2023-10-23	101.0	5400.0	0.000000	NaN	8.594339	False
4	2023-10-24	105.0	5600.0	0.039604	NaN	8.630700	False
1	2023-10-21	102.0	5200.0	0.020000	NaN	8.556606	False
2	2023-10-22	101.0	5200.0	-0.009804	NaN	8.556606	False
3	2023-10-23	101.0	5400.0	0.000000	NaN	8.594339	False
4	2023-10-24	105.0	5600.0	0.039604	NaN	8.630700	False
5	2023-10-25	107.0	5600.0	0.019048	NaN	8.630700	False
6	2023-10-26	106.0	5800.0	-0.009346	NaN	8.665786	False
7	2023-10-27	108.0	6000.0	0.018868	0.080000	8.699681	False
8	2023-10-28	110.0	6200.0	0.018519	0.078431	8.732466	False
9	2023-10-29	109.0	6400.0	-0.009091	0.079208	8.764210	False
2	2023-10-22	101.0	5200.0	-0.009804	NaN	8.556606	False
3	2023-10-23	101.0	5400.0	0.000000	NaN	8.594339	False
4	2023-10-24	105.0	5600.0	0.039604	NaN	8.630700	False
5	2023-10-25	107.0	5600.0	0.019048	NaN	8.630700	False
6	2023-10-26	106.0	5800.0	-0.009346	NaN	8.665786	False
7	2023-10-27	108.0	6000.0	0.018868	0.080000	8.699681	False
8	2023-10-28	110.0	6200.0	0.018519	0.078431	8.732466	False
9	2023-10-29	109.0	6400.0	-0.009091	0.079208	8.764210	False
5	2023-10-25	107.0	5600.0	0.019048	NaN	8.630700	False
6	2023-10-26	106.0	5800.0	-0.009346	NaN	8.665786	False
7	2023-10-27	108.0	6000.0	0.018868	0.080000	8.699681	False
8	2023-10-28	110.0	6200.0	0.018519	0.078431	8.732466	False
9	2023-10-29	109.0	6400.0	-0.009091	0.079208	8.764210	False
8	2023-10-28	110.0	6200.0	0.018519	0.078431	8.732466	False
9	2023-10-29	109.0	6400.0	-0.009091	0.079208	8.764210	False

## Observation:

This deals with preprocessing financial or stock market data for predictive modeling or analysis. Missing values in key columns are filled with appropriate statistics, such as the mean or median. Lag features are created to capture temporal dependencies, and numeric columns like trading volume are normalized to bring them onto a comparable scale. Outliers are detected and handled to prevent them from skewing model predictions. This ensures the dataset is

consistent, structured, and suitable for time-series forecasting or machine learning models.

## Task-3:

### Prompt:

Write Python code to clean IoT sensor data by forward-filling missing values, applying a rolling mean to remove drift, normalizing temperature and humidity with standard scaling, and encoding categorical sensor IDs.

### Code:

```
1  import pandas as pd
2  import numpy as np
3
4
5  def preprocess_iot_data(df: pd.DataFrame) -> pd.DataFrame:
6      """
7      Cleans and preprocesses IoT sensor data with temperature and humidity logs.
8      Handles missing values, removes sensor drift, normalizes readings, and encodes categorical sensor IDs.
9      """
10     cleaned_df = df.copy()
11
12     # 1. Handle missing values using forward fill
13     cleaned_df[["temperature", "humidity"]] = cleaned_df[
14         ["temperature", "humidity"]
15     ].fillna(method="ffill")
16
17     # 2. Remove sensor drift using rolling mean (window=3)
18     cleaned_df["temperature"] = (
19         cleaned_df["temperature"].rolling(window=3, min_periods=1).mean()
20     )
21     cleaned_df["humidity"] = (
22         cleaned_df["humidity"].rolling(window=3, min_periods=1).mean()
23     )
24
25     # 3. Normalize readings using standard scaling: (x - mean)/std
26     for col in ["temperature", "humidity"]:
27         mean_val = cleaned_df[col].mean()
28         std_val = cleaned_df[col].std()
29         cleaned_df[col] = (cleaned_df[col] - mean_val) / std_val
30
```

```

31     # 4. Encode categorical sensor IDs as integers
32     cleaned_df["sensor_id_encoded"] = (
33         cleaned_df["sensor_id"].astype("category").cat.codes
34     )
35
36     return cleaned_df.reset_index(drop=True)
37
38
39 # --- Example Usage ---
40 if __name__ == "__main__":
41     # Sample IoT sensor data
42     raw_data = {
43         "sensor_id": ["S1", "S2", "S1", "S3", "S2", "S1", "S3", "S2"],
44         "timestamp": [
45             "2023-10-26 08:00:00",
46             "2023-10-26 08:05:00",
47             "2023-10-26 08:10:00",
48             "2023-10-26 08:15:00",
49             "2023-10-26 08:20:00",
50             "2023-10-26 08:25:00",
51             "2023-10-26 08:30:00",
52             "2023-10-26 08:35:00",
53         ],
54         "temperature": [22.5, 23.0, None, 24.1, 23.5, 22.8, 24.0, None],
55         "humidity": [45, None, 47, 50, 48, 46, None, 49],
56     }
57
58     raw_df = pd.DataFrame(raw_data)
59

```

```

50         "2023-10-26 08:25:00",
51         "2023-10-26 08:30:00",
52         "2023-10-26 08:35:00",
53     ],
54     "temperature": [22.5, 23.0, None, 24.1, 23.5, 22.8, 24.0, None],
55     "humidity": [45, None, 47, 50, 48, 46, None, 49],
56 }
57
58 raw_df = pd.DataFrame(raw_data)
59
60 print("--- Original IoT Dataset ---")
61 print(raw_df)
62 print("\n" + "=" * 50 + "\n")
63
64 # Preprocess the dataset
65 cleaned_df = preprocess_iot_data(raw_df)
66
67 print("--- Cleaned IoT Dataset ---")
68 print(cleaned_df)
69

```

## Output:

```
C:\Users\venub\OneDrive\Desktop\AIAC_Lab\Lab-17>python 17_1.py
```

```
--- Original IoT Dataset ---
```

	sensor_id	timestamp	temperature	humidity
0	S1	2023-10-26 08:00:00	22.5	45.0
1	S2	2023-10-26 08:05:00	23.0	NaN
2	S1	2023-10-26 08:10:00	NaN	47.0
3	S3	2023-10-26 08:15:00	24.1	50.0
4	S2	2023-10-26 08:20:00	23.5	48.0
5	S1	2023-10-26 08:25:00	22.8	46.0
6	S3	2023-10-26 08:30:00	24.0	NaN
7	S2	2023-10-26 08:35:00	NaN	49.0

```
=====
```

--- Cleaned IoT Dataset ---

	sensor_id	timestamp	temperature	humidity	sensor_id_encoded
0	S1	2023-10-26 08:00:00	-1.622195	-1.259469	0
1	S2	2023-10-26 08:05:00	-1.030513	-1.259469	1
2	S1	2023-10-26 08:10:00	-0.833285	-0.742764	0
3	S3	2023-10-26 08:15:00	0.428969	0.549000	2
4	S2	2023-10-26 08:20:00	0.823424	1.324058	1
5	S1	2023-10-26 08:25:00	0.665642	1.065705	0
6	S3	2023-10-26 08:30:00	0.586751	0.032294	2
7	S2	2023-10-26 08:35:00	0.981206	0.290647	1
0	S1	2023-10-26 08:00:00	-1.622195	-1.259469	0
1	S2	2023-10-26 08:05:00	-1.030513	-1.259469	1
2	S1	2023-10-26 08:10:00	-0.833285	-0.742764	0
3	S3	2023-10-26 08:15:00	0.428969	0.549000	2
4	S2	2023-10-26 08:20:00	0.823424	1.324058	1
5	S1	2023-10-26 08:25:00	0.665642	1.065705	0
6	S3	2023-10-26 08:30:00	0.586751	0.032294	2
7	S2	2023-10-26 08:35:00	0.981206	0.290647	1
1	S2	2023-10-26 08:05:00	-1.030513	-1.259469	1
2	S1	2023-10-26 08:10:00	-0.833285	-0.742764	0
3	S3	2023-10-26 08:15:00	0.428969	0.549000	2
4	S2	2023-10-26 08:20:00	0.823424	1.324058	1
5	S1	2023-10-26 08:25:00	0.665642	1.065705	0
6	S3	2023-10-26 08:30:00	0.586751	0.032294	2
7	S2	2023-10-26 08:35:00	0.981206	0.290647	1
2	S1	2023-10-26 08:10:00	-0.833285	-0.742764	0
3	S3	2023-10-26 08:15:00	0.428969	0.549000	2
4	S2	2023-10-26 08:20:00	0.823424	1.324058	1
5	S1	2023-10-26 08:25:00	0.665642	1.065705	0
6	S3	2023-10-26 08:30:00	0.586751	0.032294	2
7	S2	2023-10-26 08:35:00	0.981206	0.290647	1
4	S2	2023-10-26 08:20:00	0.823424	1.324058	1
5	S1	2023-10-26 08:25:00	0.665642	1.065705	0
6	S3	2023-10-26 08:30:00	0.586751	0.032294	2
7	S2	2023-10-26 08:35:00	0.981206	0.290647	1
5	S1	2023-10-26 08:25:00	0.665642	1.065705	0
6	S3	2023-10-26 08:30:00	0.586751	0.032294	2
7	S2	2023-10-26 08:35:00	0.981206	0.290647	1
7	S2	2023-10-26 08:35:00	0.981206	0.290647	1

C:\Users\venub\OneDrive\Desktop\AIAC\_Lab\Lab-17>

## Observation:

This task involves preprocessing IoT or sensor data, which often contains missing readings, noise, and inconsistencies due to sensor drift or device errors. Missing values are filled using appropriate imputation methods, and rolling averages or other smoothing

techniques are applied to reduce noise. Sensor IDs or categorical features are encoded to numeric form so that models can process them effectively. The cleaned and normalized dataset is then ready for tasks like anomaly detection, predictive maintenance, or trend analysis.

## Task-4:

### Prompt:

The prompt asked to clean reviews by lowercasing text, removing HTML tags, encoding using TF-IDF or embeddings, filling missing ratings with the median, normalizing ratings (0–10 → 0–1), and creating a before–after summary.

### Code:

```
1 import pandas as pd
2 import re
3
4 # Basic stopwords list
5 STOPWORDS = set(
6     [
7         "i",
8         "me",
9         "my",
10        "we",
11        "our",
12        "you",
13        "your",
14        "he",
15        "she",
16        "it",
17        "they",
18        "them",
19        "this",
20        "that",
21        "a",
22        "an",
23        "the",
24        "and",
25        "or",
26        "but",
```

```
27     "if",
28     "to",
29     "of",
30     "in",
31     "on",
32     "for",
33     "with",
34     "as",
35     "is",
36     "are",
37     "was",
38     "were",
39     "be",
40     "been",
41     "has",
42     "have",
43     "had",
44     "do",
45     "does",
46     "did",
47     "not",
48 ]
49 )
50
51
52 def clean_review text(text: str) -> str:
```

```

52 def clean_review_text(text: str) -> str:
53     """Cleans review text by removing punctuation, lowercasing, and removing s
54     if not isinstance(text, str):
55         return ""
56     # Lowercase
57     text = text.lower()
58     # Remove URLs
59     text = re.sub(r"http\S+|www\S+", "", text)
60     # Remove punctuation and special characters
61     text = re.sub(r"^[a-z\s]", "", text)
62     # Remove stopwords
63     words = [word for word in text.split() if word not in STOPWORDS]
64     return " ".join(words)
65
66
67 def clean_movie_reviews(df: pd.DataFrame) -> pd.DataFrame:
68     """
69     Cleans a movie reviews dataset:
70     - Removes duplicates
71     - Cleans review text
72     - Handles missing ratings
73     """
74     cleaned_df = df.copy()
75
76     # 1. Remove duplicates
77     cleaned_df.drop_duplicates(subset=["review text"], keep="first", inplace=True)

```



```

77     cleaned_df.drop_duplicates(subset=["review_text"], keep="first", inplace=True)
78
79     # 2. Fill missing ratings with median
80     if "rating" in cleaned_df.columns:
81         cleaned_df["rating"] = pd.to_numeric(cleaned_df["rating"], errors="coerce")
82         median_rating = cleaned_df["rating"].median()
83         cleaned_df["rating"].fillna(median_rating, inplace=True)
84         cleaned_df["rating"] = cleaned_df["rating"].astype(int)
85
86     # 3. Clean review text
87     cleaned_df["cleaned_review"] = cleaned_df["review_text"].apply(clean_review_text)
88
89     # 4. Reset index
90     return cleaned_df.reset_index(drop=True)
91
92
93 # --- Example Usage ---
94 if __name__ == "__main__":
95     raw_data = {
96         "review_id": [101, 102, 103, 104, 105, 106],
97         "review_text": [
98             "I loved this movie! It was fantastic. http://example.com",
99             "Terrible movie... would not recommend! #fail",
100             "I loved this movie! It was fantastic.", # Duplicate
101             "Average movie, some good parts, some bad."

```

ROBLEMS OUTPUT DEBUG CONSOLE PORTS AZURE QUERY RESULTS POSTGRESQL QUERY RESULTS TERMINAL

```

94 if __name__ == "__main__":
95     raw_data = {
96         "review_id": [101, 102, 103, 104, 105, 106],
97         "review_text": [
98             "I loved this movie! It was fantastic. http://example.com",
99             "Terrible movie... would not recommend! #fail",
100             "I loved this movie! It was fantastic.", # Duplicate
101             "Average movie, some good parts, some bad.",
102             None,
103             "Best movie ever! A must-watch!",
104         ],
105         "rating": [5, 1, 5, None, 3, 5],
106     }
107
108     raw_df = pd.DataFrame(raw_data)
109
110     print("--- Original Movie Reviews Dataset ---")
111     print(raw_df)
112     print("\n" + "=" * 50 + "\n")
113
114     # Clean dataset
115     cleaned_df = clean_movie_reviews(raw_df)
116
117     print("--- Cleaned Movie Reviews Dataset ---")
118     print(cleaned_df)
119

```

## Output:

```

C:\Users\venub\OneDrive\Desktop\AIAC_Lab\Lab-17>python 17_1.py
--- Original Movie Reviews Dataset ---

```

	review_id	review_text	rating
0	101	I loved this movie! It was fantastic. <a href="http://e...">http://e...</a>	5.0
1	102	Terrible movie... would not recommend! #fail	1.0
2	103	I loved this movie! It was fantastic.	5.0
3	104	Average movie, some good parts, some bad.	NaN
4	105	None	3.0
5	106	Best movie ever! A must-watch!	5.0

```

=====

```

```
--- Cleaned Movie Reviews Dataset ---
  review_id  ...                               cleaned_review
0         101  ...                               loved movie fantastic
1         102  ...      terrible movie would recommend fail
2         103  ...                               loved movie fantastic
3         104  ...  average movie some good parts some bad
4         105  ...
5         106  ...          best movie ever mustwatch
```

```
[6 rows x 4 columns]
```

```
C:\Users\venub\OneDrive\Desktop\AIAC_Lab\Lab-17>
```

## Observation:

This focuses on cleaning and preparing textual review data, such as movie or product reviews, for sentiment analysis or recommendation systems. Text preprocessing includes lowercasing, removing HTML tags, and optionally removing stopwords or special characters. The text is then encoded using methods like TF-IDF or embeddings to transform it into numeric features usable by machine learning models. Ratings are preprocessed by filling missing values with the median and normalizing them to a 0–1 scale. The task may also include generating before-and-after summaries to compare the raw and cleaned datasets, ensuring the text is ready for further analysis or modeling.